

Balancing an Inverted Double Pendulum with Reinforcement Learning

Cem Alptürk¹ Pukashawar Pannu²

¹ce5368al-s@student.lu.se ²pu6218pa-s@student.lu.se

Abstract: Controlling an inverted double pendulum on a cart is an inherently difficult problem. The system is chaotic meaning that the dynamics cannot be predicted over a longer time span. A way to control such a system is to use the current state to determine an action that moves the cart appropriately to balance the double pendulum. One such way is by reinforcement learning. The approach in this project uses a model free controller trained with Deep Q-learning to control the double pendulum. Benefits of Deep Q-learning over regular Q-learning is that the controller can be used for situations that are not presented during training. This has been confirmed in the results for the single pendulum model, an easier problem to control that was used to confirm the implementation of the learning algorithm. The double pendulum is more difficult with its chaotic nature and the limitation that the cart with a single degree of freedom is trying to control multiple states. Results show that the double pendulum can be balanced for up-to three seconds with a network of 10 layers and 30 nodes in each layer. The agent was trained for a limited number of episodes limiting the amount of experience and knowledge it could gather about the problem. We believe that by tuning hyperparameters, network size and training one could train an agent to balance the double pendulum for far longer using Deep Q-Learning.

1. Introduction

The goal of this project is to teach an agent to control a double pendulum on a cart to be standing upright, in its equilibrium point, using reinforcement learning. A double pendulum is a highly chaotic system, where its behavior heavily depends on its initial conditions. Controlling the system with a cart limited to a single degree of freedom is a difficult but solvable problem [6]. The aim is to create an agent that can learn from its own experiences to control the double pendulum, by applying a force on the cart in the horizontal direction, such that it stays upright as seen in Figure 1. Although this is an academic problem, the results could be applied to any type of learning problem since the agent is unaware of what it is working on [5].




Figure 1. A double pendulum being controlled by a trained agent that is applying $-1N$ in the horizontal direction.

1.1 Reinforcement Learning

The main idea of reinforcement learning is to learn to make decisions based on past experiences. In this case this means to

train an agent to perform the best action possible in order to control the pendulums. For each action taken by the agent, the environment returns a new state and a reward for that action. The reward is calculated based on how close the system is to its equilibrium point. The goal of training the agent is to optimize gained reward for each action taken.

2. Modeling

The system consists of a cart that can move in the horizontal direction and two pendulums that are coupled to each other such that there is no collision between any parts. Since this project is not about modeling the double pendulum, but rather about controlling it, the equations of motion have been borrowed from [7]. The states of the system are represented by,

$$x := [y \quad \dot{y}]^T, \quad y := [q \quad \theta_1 \quad \theta_2]^T \quad (1)$$

- q is the position of the cart on the horizontal axis.
- θ_1 is the angle of the inner pendulum with respect to the vertical axis in the clockwise direction.
- θ_2 is the angle of the outer pendulum with respect to the vertical axis in the clockwise direction.
- \dot{q} is the velocity of the cart in the horizontal axis.
- $\dot{\theta}_1$ is the angular velocity of the inner pendulum.
- $\dot{\theta}_2$ is the angular velocity of the outer pendulum.

The system parameters are represented by,

- $m = 1[kg]$ is the mass of the cart.
- $m_1 = 0.1[kg]$ is the mass of the inner pendulum.
- $m_2 = 0.1[kg]$ is the mass of the outer pendulum.
- $l_1 = 1[m]$ is the length of the inner pendulum.
- $l_2 = 1[m]$ is the length of the outer pendulum.
- $g = 9.81[m/s^2]$ is the gravitational acceleration.

The equations of motions for this system [7], are adjusted such there is no friction or disturbances. The input to the system is a horizontal force that is acting on the cart, represented by u [N].

$$M(y)\ddot{y} = f(y, \dot{y}, u) \quad (2)$$

$$M(y) := \begin{bmatrix} m + m_1 + m_2 & l_1(m_1 + m_2) \cos \theta_1 & m_2 l_2 \cos \theta_2 \\ l_1(m_1 + m_2) \cos \theta_1 & l_1^2(m_1 + m_2) & l_1 l_2 m_2 \cos(\theta_1 - \theta_2) \\ l_2 m_2 \cos \theta_2 & l_1 l_2 m_2 \cos(\theta_1 - \theta_2) & l_2^2 m_2 \end{bmatrix} \quad (3)$$

$$f(y, \dot{y}, u) := \begin{bmatrix} l_1(m_1 + m_2)(\dot{\theta}_1)^2 \sin \theta_1 + m_2 l_2 (\dot{\theta}_2)^2 \sin \theta_2 + u \\ -l_1 l_2 m_2 (\dot{\theta}_2)^2 \sin(\theta_1 - \theta_2) + g(m_1 + m_2) l_1 \sin \theta_1 \\ l_1 l_2 m_2 (\dot{\theta}_1)^2 \sin(\theta_1 - \theta_2) + g l_2 m_2 \sin \theta_2 \end{bmatrix} \quad (4)$$

3. System Design

A couple of reinforcement learning methods are Q-learning and Deep Q-learning.

3.1 Q-learning

Q-learning is a model free method [8] where the agent has no information about the system it is trying to control, thus it can be said that it is a black box model. The agent stores all its experiences in a Q-table that contains the quality of a state-action combination.

$$Q : S \times A \rightarrow \mathbb{R} \quad (5)$$

$Q(s, a) :=$ Quality of action a given state s

A trained agent decides on the action to perform by looking up the best action a given state s from the Q-table. In the case where multiple actions give the same values in the Q-table, the action is selected randomly among those. The selected action is performed on the environment, which is the inverted pendulum, and a reward is calculated based on the new state of the system. In the case where there is no previous knowledge about the current state in the Q-table, the action is selected randomly. Essentially the agent tries to find the action that will give the best reward.

$$a^* = \max_i Q(s, a_i) \quad (6)$$

The agent uses an ϵ -greedy policy to decide on the next move, where the agent performs a random action with probability ϵ or performs the best action based on its current policy with probability $1 - \epsilon$. Having a high ϵ value will correspond to performing randomly and exploring the state-action space but will delay the convergence of the Q-table. Having a low ϵ value

can cause the agent to rarely try new things but instead will try to exploit its current solution to the problem. A balance between these problems can be obtained by starting with a high ϵ value and decaying it over time to allow the agent to try new things in the beginning and slowly exploit the best solution it has found.

Each state-action pair in a trained Q-table, represents the expected value in Equation (7). These values represent the expected return of rewards for future actions if action a_t is taken at state s_t .

$$Q(s_t, a_t) = E[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | s_t, a_t] \quad (7)$$

Training the agent corresponds to populating and updating the Q-table based on past experiences of the agent. After performing an action given the current state of the system, the agent receives a reward from the environment and the values in the Q-table are updated according to the state-action combination and the reward received. Equation (8) is known as the Bellman equation.

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha)Q(s_t, a_t) + \alpha(r_t + \gamma \max_i Q(s_{t+1}, a_i)) \quad (8)$$

$$0 < \alpha \leq 1, \quad 0 \leq \gamma \leq 1$$

After performing action a_t given the state s_t , the agent receives the new state s_{t+1} and the reward r_t . The value $Q(s_t, a_t)$ is updated by Equation (8). α is the learning rate and γ is the discount factor. Since the reward for a certain action might not come instantly, the discount factor tries to compensate for a delayed reward by estimating the reward that the agent will receive if it takes the best action a_{t+1} for the state s_{t+1} . Low values of γ will cause the agent to only consider near future rewards, while a γ value close to 1 will consider future rewards.

This method requires a discrete set of actions and states. For the double pendulum problem, this can be achieved by discretizing the states and actions. However since this is a chaotic system, it would require a fine discretization resulting in a massive Q-table. Although this can work in theory, there are practical implementation limitations such as limited memory size. This problem can be resolved by using the Deep Q-learning method.

3.2 Deep Q-learning

In Deep Q-learning, the Q-table is replaced with a deep neural network [5] that tries to learn Q-values instead of memorizing them. A benefit of using Deep Q-learning is that it uses less memory and it has more predictive power on unseen states. The network takes in the system states as input arguments and predicts the Q-values.

$$NN_Q : \mathbb{R}^m \rightarrow \mathbb{R}^n \quad (9)$$

where m is the state size and n is the size of the action space. The agent chooses the action to perform by picking the largest Q-value from the neural network prediction.

The neural network is trained by a method called experience replay that is based on simulating the system and recording all states, rewards and actions taken in memory [5]. During




Figure 2. Illustration of a neural network with 3 state inputs and 4 Q-value outputs that can be mapped to actions.

experience replay, a random batch is taken from the memory and is used to train the network. The reason for taking random batches from memory and not sequential data, is to reduce the correlation between the data points. The simulation ends when it reaches a termination condition that is defined in the environment. For the double pendulum case the termination condition is when the outer pendulum angle becomes larger than a defined maximum value. One full simulation from start to termination is called an episode. The stored experience is used to fit the neural network once there is enough experience stored in the memory.

$$Exp \leftarrow (s_t, a_t, r_t, T_t, s_{t+1}) \quad (10)$$

Where s_t is the state at time t , a_t the action taken by the system at time t , r_t the gained reward for the taken action, T_t indicating whether the system has terminated and s_{t+1} the next state in time. The experience is used to map a target for the neural network training. Targets are Q-values for each action given a state. Since the experience only holds data about the actual action that was taken for a given state the other Q-values for the same state have to be estimated to complete the entire target vector. There are several ways to estimate the unknown Q-values for a given state. Although the NN_Q network can be used to predict the target values for the other actions it causes the targets to be non-stationary leading to the network chasing its own tail [4]. A separate neural network, NN_{target} , was used to avoid the non-stationary targets problem. NN_{target} has the same network architecture as NN_Q . After a number of episodes the weights of NN_Q are mapped to NN_{target} to give better estimations of the unknown Q-values.

The actual value y_m replaces the corresponding predicted Q-value from $NN_{target}(s_m)$. Since we only have the reward for one action in a given state, we replace the targets for the untaken actions with a prediction from the Q-network. This will result in a loss of 0 for these untaken actions, thus not affecting the training. This way the network is trained only by the state-action pair that has been taken. This will however cause the implementation to be slower since a prediction has to be made before each weight update. This was solved by implementing a custom loss function in Keras where the loss only depends on the Q-value of the taken action.

Algorithm 1 Deep Q-Learning

```

for episode [0, max episodes] do
  while  $t < t_{max}$  and  $T_t$  is false do    ▷ Simulate system
     $q_{values} \leftarrow NN_Q(s_t)$ 
     $a \leftarrow \text{argmax}(q_{values})$                 ▷  $a$  is index
     $s_{t+1}, r_t, T_t \leftarrow \text{Step}(\text{ActionSpace}[a])$ 
     $Exp \leftarrow (s_t, a_t, r_t, T_t, s_{t+1})$ 
     $(s, a, r, T, s_{t+1}) \leftarrow Exp$     ▷ Exp minibatch

     $y_m = \begin{cases} r_m \\ r_m + \gamma \max(NN_{target}(s_{t+1,m})) \end{cases}$ 
     $Loss \leftarrow (y_m - NN_Q(s_{t+1,m})_{a_m})^2$ 
    Perform gradient descent on Loss
  end while
end for

```

4. Implementation

The project was implemented in Python. A separate class was created for each component Agent, Environment, Simulator and Controller. The goal of the project is to produce a Controller object that can control the Simulator. Environment class contains the problem specifics such as reward and termination functions, and links the Agent with the Simulator containing the physical model dynamics. Agent contains the learning algorithm that trains and produces the Controller. The Controller is a simple object that produces the force to apply given the current state of the system.

$$C := \mathbb{R}^M \rightarrow \mathbb{R} \quad (11)$$

All components are independent of each other making the structure modular. By using this structure it is easy to change for example the learning algorithm without having to modify anything else.

4.1 Simulator

The Simulator is responsible for the physical dynamics of the problem and for simulating the problem for given external inputs and time. Equations are on the form of RHS-equations and are solved using odeint function in Scipy.integrate package. For the double pendulum case, the simulator solves Equation (2) for \ddot{y} to produce the next state using the ODE solver.

4.2 Environment

The reward function and termination condition is implemented in the Environment class. It acts as a connector between the Agent and the Simulator. This way when the reward or the termination needs to be changed it can be done without altering the learning algorithm (Agent) or the physical model (Simulator).

4.3 Evaluation

Controller is evaluated after a set of episodes in order to keep track of the controller performance. Performance is measured by averaging the total reward for N_{eval} simulations of the

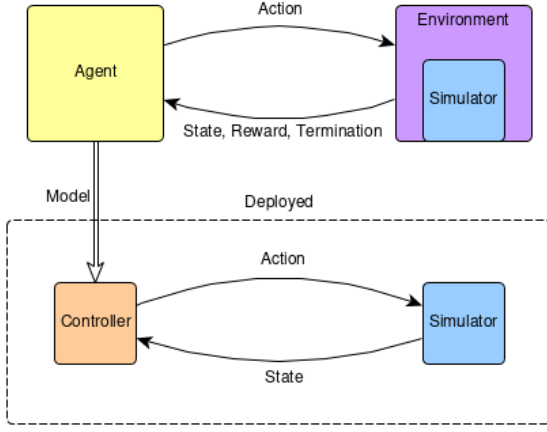


Figure 3. Code architecture.

dynamic system controlled by the Controller.

$$\frac{1}{N_{eval}} \sum_i^{N_{eval}} \sum_j^{M_i} R(s_j) \quad (12)$$

Notice that the number of simulation steps M_i for each simulation varies as a simulation may reach a termination condition sooner or later than another simulation. A satisfactory evaluation value can be used as a termination condition for the training to avoid overtraining. This mechanism has not been implemented. The training will continue until the max iterations have been reached.

5. Results

A simpler problem was used to confirm the implementation. The simpler problem being to balance a single inverted pendulum on a cart. The implementation architecture was used to simply replace the environment with one suitable for the single pendulum problem.

5.1 Balancing Single Inverted Pendulum

The single pendulum problem is non-chaotic. Equations from [1] were used to model the problem. System was terminated when the pendulum angle was greater than ± 10 degrees. Equation (13) states the reward function.

$$R(s) = (\theta_{max} - |\theta|)/\theta_{max} \quad (13)$$

where $\theta_{max} = 10$ degrees. The value was chosen by trial and error, where 10 degrees gave the best overall performance.

Table 1 shows the neural network setup for the single pendulum problem. The weights of the network were initialized with uniform HE-initialization, which is an optimized weight initialization method when using ReLU activation functions [3].

Layer	Number of Neurons	Activation
Input	4	-
Hidden 1	20	ReLU
Hidden 2	40	ReLU
Output	3	Linear

Table 1. Network parameters for the single pendulum Q-network.

Table 2 shows the training parameters for the single pendulum problem.

Max Angle	10
Step Size	0.02
Action Space	$[-10, 0, 10]$
Learning Rate	0.001
Memory	2000
Batch Size	32
Discount	0.9

Table 2. Parameters for the Deep Q-learning algorithm on the single pendulum problem.

Figure 4 shows the agent performance over episodes. The agent starts to learn how to move the cart to increase the score after a few episodes. Eventually the reward gain stops increasing. The system was simulated with the maximum 200 steps for each episode, capping the max possible reward to 200. However, 200 is only possible if the initial condition starts in the equilibrium point. The initial condition for each episode is randomized preventing the system from reaching 200 score in the evaluation.




Figure 4. Evaluated controller score for every 10 episodes. The evaluation is measured as the average score of 10 full controlled simulations. Maximum number of steps for each evaluation iteration was 200, capping the maximum possible reward to 200.

The trained controller was used to control the single pendulum for initial conditions outside the trained domain, and the controller managed to bring back the pendulum to the known domain and balance it. This can be seen in Figure 5. Results indicate that the network has managed to generalize the problem. This would not be possible with a regular Q-table approach since the agent has never seen the states outside the training domain and would end up performing random actions. Although the agent was able to bring the pendulum back to its equilibrium point, it later starts to move to the right with a small angle. The reason for this might be that the cart ended up far away from the origin, which confused the agent. This can be resolved by training the agent for more episodes.

Figure 6 and 7 shows the average Q-values during training for $\gamma = 0.1$ and $\gamma = 0.9$ respectively. It can be seen that the Q-values for a low γ value converges much faster while for

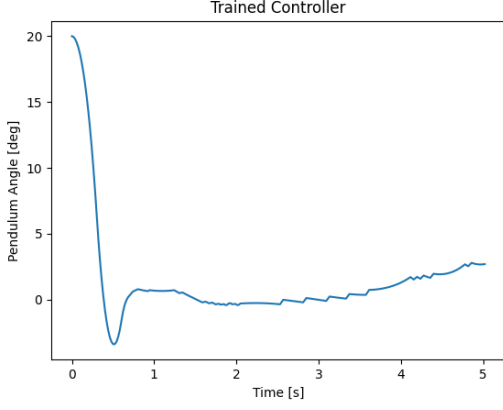


Figure 5. Changes in pendulum angle θ over time for controlled simulation of single pendulum on cart problem. The initial condition was $\theta = 20$ degrees.

the larger γ value, convergence has not happened at the end of the training.




Figure 6. Q-values for $\gamma = 0.1$.

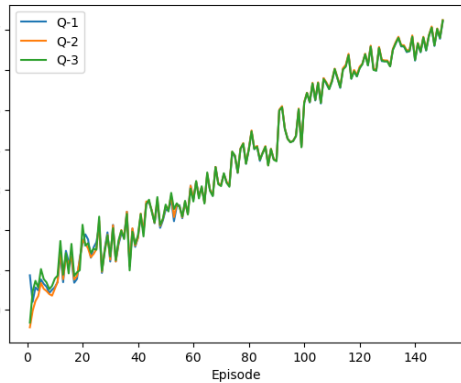


Figure 7. Q-values for $\gamma = 0.9$.

Figure 8 shows the training result for $\gamma = 0.1$, while the other hyperparameters remained the same. It can be seen that the score converges faster since the Q-values converge faster. This is due to the algorithm only taking the instant and near

future rewards in to consideration, thus being short sighted. However the overall performance showed that, with $\gamma = 0.1$ the agent was unable to generalize and even had issues with initial conditions that were close to θ_{max} . This can also be seen in the dip during the end of the training in Figure 8.

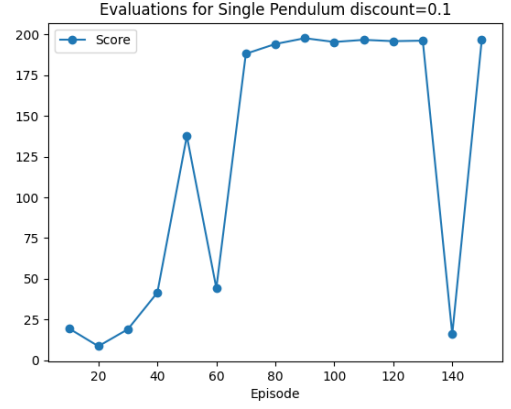


Figure 8. Evaluation of the controller with the same conditions in Table 2, except that $\gamma = 0.1$. The drop at the end may be caused due to outlier initial conditions during evaluation.

Figure 9 shows the distribution of actions taken for an agent that was trained with a larger action space. This plot shows that most of these actions were rarely used except for ± 10 . When the agent was trained for the same number of episodes, the overall performance was worse than in Figure 4. This showed that a larger action space requires more training in order to get the same performance. In this case training time can be reduced by only using the actions that were most used by the agent.




Figure 9. The distribution of used actions for an action space of $[-20, -10, -5, 0, 5, 10, 20]$. The model was trained with the same hyperparameters in Table 2.

5.2 Balancing Inverted Double Pendulum

The termination condition is given by Equation (14),

$$|\theta_1| > \theta_{max} \text{ or } |\theta_2| > \theta_{max} \quad (14)$$

where θ_1 is the angle of the inner pendulum and θ_2 is the angle of the outer pendulum. $\theta_{max} = 5$ degrees was chosen in order

to train the agent in a strict state space close to the equilibrium point. Due to the limited degrees of freedom of the system, higher values of θ_{max} can allow the pendulums to move to positions where they can not be recovered from. Reward is given by Equation (15),

$$R(s) = 1 - (\theta_1^2 + \theta_2^2) - (\dot{\theta}_1^2 + \dot{\theta}_2^2) \quad (15)$$

where the agent receives a bonus of 1 for being alive, a penalty for being further away from the equilibrium point given by $\theta_1^2 + \theta_2^2$, and a penalty for the velocity of the pendulums given by $\dot{\theta}_1^2 + \dot{\theta}_2^2$. The total score received in an episode depends on the step size of the simulation and the duration that the agent is able to survive.

Table 3 shows the neural network setup for the double pendulum problem. The weights of the network were initialized with uniform HE-initialization similar to the single pendulum problem.

Layer	Number of Neurons	Activation
Input	6	-
Hidden 1-10	30	ReLU
Output	3	Linear

Table 3. Network parameters for the double pendulum Q-network.

Table 4 shows the training parameters for the double pendulum problem.

Max Angle	5
Step Size	0.02
Action Space	[-1,0,1]
Learning Rate	0.0001
Memory	2000
Batch Size	32
Discount	0.9

Table 4. Parameters for the Deep Q-learning algorithm on the double pendulum problem.

Figure 10 shows the agents score over 500 episodes. During training the agent was able to balance the pendulum up to three seconds. A comparison between the agents performance, baseline and a random controller is also illustrated in the figure. Although the trained agent cannot balance the pendulums for a long time, it performs much better than the others.




Figure 10. Evaluated controller score for every 10 episodes. The evaluation is measured as the average score of 10 full controlled simulations. The baseline is measured by evaluating 200 simulations without a controller which lasted on average for 0.5 seconds. A controller that performs actions randomly achieves an average score of 32, estimated the same way as the baseline.

Figure 11 shows the average Q-values of the network during training. It can be seen that the Q-values have converged and can be interpreted as the network having reached its limit and training being complete.




Figure 11. Q-values during training.

5.3 Training Time

The training time of the agent heavily depends on the hyperparameters and the overall performance of the agent. In the first few episodes, due to the agents lack of experience the episodes will terminate faster, thus resulting in shorter episodes. As the agent gains more experience, it learns to survive for longer, increasing the computation time for a single episode. Table 5 and 6 show the computation and simulation time differences for a change in the termination condition.

Single Pendulum		
Episode	Computation	Simulation
1	0.72 s	0.56 s
⋮	⋮	⋮
50	24.73 s	4 s
⋮	⋮	⋮
100	28.52 s	4 s

Table 5. Computation and simulation times for episodes with hyperparameters specified in Table 2.

Single Pendulum		
Episode	Computation	Simulation
1	0.76 s	0.6 s
⋮	⋮	⋮
50	6.03 s	0.86 s
⋮	⋮	⋮
100	26.77 s	4 s

Table 6. Computation and simulation times for episodes with hyperparameters specified in Table 2, with a modification of $\theta_{max} = 20$ degrees.

Double Pendulum		
Episode	Computation	Simulation
1	2.95 s	0.8 s
⋮	⋮	⋮
300	17.38 s	1.96 s
⋮	⋮	⋮
500	29.13 s	2.62 s

Table 7. Computation and simulation times for episodes with hyperparameters specified in Table 4.

6. Discussion

The double pendulum is a highly chaotic system and with the cart limited to a single degree of freedom it is difficult to control. In situations where the agent receives two states that are very close to each other, the applied action may result in very different outcomes which makes the learning of the expected reward for that action very difficult. We have found that deeper networks perform better than wider ones for this problem. This may be due to wider networks having less non-linearity compared to deeper ones since deep networks have more coupled activations.

After playing around with the exploration rate ϵ , we have found that a dynamic exploration rate performs better for this problem, rather than a static one. Higher exploration rate forces the agent to take actions that it normally would avoid causing it to explore the action space more. For Q-learning, this is a desired thing since the purpose is to fill and update the Q-table. However a high exploration rate leads to the agent exploring things that it should avoid, rather than finding the best action to perform. In Q-learning, this is not an issue since

the Q-values of the state-action pairs do not directly affect each other. However for Deep Q-learning, every weight update affects all Q-values for the state-action pairs. In the case where the learning rate is too low, the agent will almost always prefer to take the same actions which can cause overfitting in the network. This is why we preferred to train the agent with a dynamic exploration rate, where the agent performs randomly in the beginning and over time becomes more deterministic. However we still have a minimum ϵ , to allow the agent to explore during the later episodes of training. This prevents the agent from overfitting to states close to the equilibrium point. For the same reason, we preferred to use a strict termination condition. Since the agents purpose is to control the pendulums around the equilibrium, there is little to no advantage in allowing the pendulum to move too far from the equilibrium point.

When we used a discount factor closer to 1, we found that it takes longer for the Q-values to converge, and they converged to larger values, as seen in Figure 7. Although this specific problem does not have any delayed rewards, such as reaching a goal in a game, it made more sense to use a high discount factor γ , so that the agent can take into account the consequences of its actions, such as if an action makes the pendulum overshoot the equilibrium point or not. Our results indicated that using a higher γ value helped the agent generalize much better and was more resilient to near edge starting conditions.

As seen in Table 7 tuning hyperparameters is a difficult process. Due to the time the training takes, evaluating the performance of a different set of hyperparameters is challenging. Since the training process highly depends on the hyperparameters, the time it takes for one episode can be vastly different for different hyperparameters. Aside from the time it takes to get the results, evaluating the effect of the hyperparameters is not straight forward. Even when using the same set of hyperparameters we can get different results. Trying to estimate the effect of a new hyperparameter is problematic. Due to these reasons, a trial and error approach was more effective. We believe that with the right hyperparameters, the double pendulum can be controlled for longer.

The duration that the agent survives for an episode does not directly correlate with the agents performance. Different hyperparameters, reward functions and termination conditions affect the duration the agent is able to survive, which makes evaluating hyperparameters problematic. As seen in Table 5 and 6, the agent was trained with a more flexible termination condition, which allowed it to survive longer in an episode, however it takes more episodes for the Q-values to converge since the state-action space is much larger than before which takes more training time. The simulation time for episode 50 in Table 5 cannot be directly compared with the corresponding value in Table 6.

We have tried other model-free methods such as Linear Q-learning [9] and Double Deep Q-learning [2], but they did not achieve better performance than our Deep Q-learning [5] implementation.

References

- [1] R. Florian. “Correct equations for the dynamics of the cart-pole system” (2005).
- [2] H. van Hasselt, A. Guez, and D. Silver. “Deep reinforcement learning with double q-learning”. *CoRR* **abs/1509.06461** (2015). arXiv: [1509.06461](https://arxiv.org/abs/1509.06461). URL: <http://arxiv.org/abs/1509.06461>.
- [3] K. He, X. Zhang, S. Ren, and J. Sun. “Delving deep into rectifiers: surpassing human-level performance on imagenet classification”. *IEEE International Conference on Computer Vision (ICCV 2015)* **1502** (2015). doi: [10.1109/ICCV.2015.123](https://doi.org/10.1109/ICCV.2015.123).
- [4] T. Hester, M. Vecerik, O. Pietquin, M. Lanctot, T. Schaul, B. Piot, A. Sendonaris, G. Dulac-Arnold, I. Osband, J. Agapiou, J. Leibo, and A. Gruslys. “Learning from demonstrations for real world reinforcement learning” (2017).
- [5] V. Mnih, K. Kavukcuoglu, D. Silver, A. Graves, I. Antonoglou, D. Wierstra, and M. Riedmiller. “Playing atari with deep reinforcement learning” (2013).
- [6] L. Moysis. *Balancing a double inverted pendulum using optimal control and laguerre functions*. 2016. doi: [10.13140/RG.2.1.2948.6486](https://doi.org/10.13140/RG.2.1.2948.6486).
- [7] Z. Neusser and M. Valasek. “Control of the double inverted pendulum on a cart using the natural motion”. *Acta Polytechnica* **53** (2013). doi: [10.14311/AP.2013.53.0883](https://doi.org/10.14311/AP.2013.53.0883).
- [8] C. Watkins and P. Dayan. “Technical note: q-learning”. *Machine Learning* **8** (1992), pp. 279–292. doi: [10.1007/BF00992698](https://doi.org/10.1007/BF00992698).
- [9] Y. Zheng, S. Luo, and Z. Lv. “Control double inverted pendulum by reinforcement learning with double cmac network”. In: vol. 4. 2006, pp. 639–642. doi: [10.1109/ICPR.2006.416](https://doi.org/10.1109/ICPR.2006.416).